



CI/CD



Continuous delivery continuous
integration



Summary

- Overview of CI/CD
 - Why CI/CD important?
- CI/CD for microservice
 - CD and Facebook/OANDA case studies
- CI and Uber Submit Queue pipeline
- Limitation and Future Work

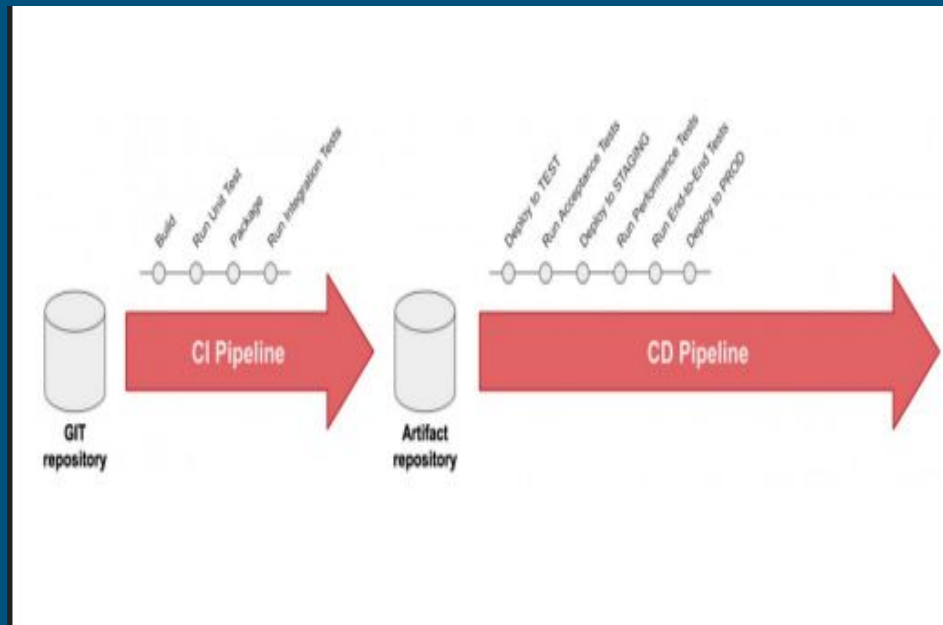
Overview of CI/CD

Continuous Integration - regularly build, test, and merge code changes into main branch.

Continuous Deployment - automatically test and release changes from the repo to production.

CI is generally a standard across all software projects.

CD on the other hand is not. Example: Aerospace Industries, Healthcare, etc.



Why is CI/CD useful?

- Release software with less risks
 - CI/CD takes care of the automated testing, deployment and rollbacks.
 - Can improve developer productivity.
 - More automation, less context switching, etc
 - Ship features and fix bugs faster.
 - Ability to push out small changes and iterate on them continually.
-

CI/CD for microservices

- Requirements

- Highly cohesive and loosely coupled services and teams
 - Teams and engineers that can make key-decisions independently

- Challenges:

- Tools to support deployment and experience in managing them.
 - Senior Management buy-in.
 - Investment on the tools

Despite the challenges, a lot of companies use CI/CD to manage their release.

CI/CD in practice

Steps involved:

- Code Review
- Testing
 - Unit tests, integration tests, performance tests, etc
- Release engineering
 - Assess the risk and manage the deployment
- Deployment

Continuous Deployment

- Blue-green deployment
 - Deploy to a small fraction of people and dial it up
- Dark launches
 - Launch during non
 - peak hours
- Staging
 - Test the builds in multiple staging environments
 - Simulate some traffic

Transition to CI/CD

- Automated Testing infrastructure
 - Unit tests, Integration tests, Shadow tests, performance tests, etc
- Deployment Management System
 - Code reviews, VCS, deployment scheduling, staging pipelines, Rollbacks

CD Case Studies At Facebook and OANDA

- OANDA

- Currency trading system that manages trades worth many billions every day
- Small team with about 100 engineers
- Code check-in -> Deployment Pipelines -> ad-hoc alert system using emails

- Facebook

- Billions of queries per second
- 1000s of engineers
- Code check-in -> Release engineering -> Error reporting (SEV)

Goal: Using CI to Keep the Master Branch Green

- Instantly release new features from any commit point in the mainline
- Can roll back to any previously committed change
- Improve engineer's productivity

Idea behind:

- Speculatively build possible outcomes for pending changes.
- Try to minimize the number of builds and parallelly build independent changes.

Limitation and Future Work

- No order for non-independent changes: Possible starvation for small changes
 - Possible to implement a better machine learning method to train the prediction model
- Should abort a build which is near its completion but likelihood of success drops?
- Batch independent changes